


What is S3?

S3 stand for Simple Storage Service.

It is a software offering by Amazon Web Services (AWS) that essentially acts like **your own personal storage space in the cloud**.

You can think of S3 the same way you would an external hard drive.





What is the application for S3?

Like hard disk storage, there are numerous applications for S3.

S3 offers file versioning (i.e. allows for multiple versions of a file to exist in a storage location, also referred to as buckets)

You can scale your storage, in other words it's not a fixed size like an external hard drive would be.

It's cost effective. Depending on how frequent you will need to access your files, you can choose various types of storage types and further reduce the amount you pay.


There are analytics use cases that will call for S3 usage and you can even use it to host simple HTML sites!

What is an S3 bucket?

An S3 bucket is a storage unit.

It can contain **folders and files**, also referred to as **objects**

S3 bucket names are unique, so you can't share the same bucket name with anyone, even if they have a different AWS account.



STORAGE

Free Tier 12 MONTHS FREE

Amazon S3

5 GB

of standard storage

Secure, durable, and scalable object storage infrastructure.

5 GB of Standard Storage

20,000 Get Requests

2,000 Put Requests

^

Create a Bucket

```
import boto
```

```
def create_bucket(bucket_name):
```

```
    """
```

Create a bucket. If the bucket already exists and you have access to it, no error will be returned by AWS.

Note that bucket names are global to S3 so you need to choose a unique name.

```
    """
```

```
    s3 = boto.connect_s3()
```

```
    # First let's see if we already have a bucket of this name.
```

```
    # The lookup method will return a Bucket object if the
```

```
    # bucket exists and we have access to it or None.
```

```
    bucket = s3.lookup(bucket_name)
```

```
    if bucket:
```

```
        print 'Bucket (%s) already exists' % bucket_name
```

```
else:
    # Let's try to create the bucket. This will fail if
    # the bucket has already been created by someone else.
    try:
        bucket = s3.create_bucket(bucket_name)
    except s3.provider.storage_create_error, e:
        print 'Bucket (%s) is owned by another user' %
bucket_name
    return bucket
```

Create a Bucket in a Specific Location

```
import boto
```

```
from boto.s3.connection import Location
```

```
def create_bucket(bucket_name, location=Location.DEFAULT):  
    """
```

Create a bucket. If the bucket already exists and you have access to it, no error will be returned by AWS.

Note that bucket names are global to a S3 region or location so you need to choose a unique name.

bucket_name - The name of the bucket to be created.


```

location - The location in which the bucket should be
           created. The Location class is a simple
           enum-like static class that has the following attributes:
           DEFAULT|EU|USWest|APNortheast|APSoutheast
"""
s3 = boto.connect_s3()
# First let's see if we already have a bucket of this name.
# The lookup method will return a Bucket object if the
# bucket exists and we have access to it or None.
bucket = s3.lookup(bucket_name)
if bucket:
    print 'Bucket (%s) already exists' % bucket_name
else:
    # Let's try to create the bucket. This will fail if
    # the bucket has already been created by someone else.
    try:
        bucket = s3.create_bucket(bucket_name,
location=location)
    except s3.provider.storage_create_error, e:
        print 'Bucket (%s) is owned by another user' %
bucket_name
    return bucket

```

Store Private Data

```
import boto
```

```
def store_private_data(bucket_name, key_name, path_to_file):
```

```
    """
```

Write the contents of a local file to S3 and also store custom metadata with the object.

bucket_name The name of the S3 Bucket.

key_name The name of the object containing the data in S3.

path_to_file Fully qualified path to local file.

```
    """
```

```
s3 = boto.connect_s3()
```

```
bucket = s3.lookup(bucket_name)
```

Get a new, blank Key object from the bucket. This Key object only

exists locally until we actually store data in it.

key = bucket.new_key(key_name)

First let's demonstrate how to write string data to the Key

data = 'This is the content of my key'

key.set_contents_from_string(data)

Now fetch the data from S3 and compare

stored_key = bucket.lookup(key_name)

stored_data = stored_key.get_contents_as_string()

assert stored_data == data

Now, overwrite the data with the contents of the file

key.set_contents_from_filename(path_to_file)

return key

Store Metadata with an Object

```
import boto
```

```
def store_metadata_with_key(bucket_name,  
                             key_name,  
                             path_to_file,  
                             metadata):
```

```
    """
```

Write the contents of a local file to S3 and also store custom metadata with the object.

bucket_name The name of the S3 Bucket.

key_name The name of the object containing the data in S3.

path_to_file Fully qualified path to local file.

metadata A Python dict object containing key/value data you would like associated with the object.

For example: {'key1':'value1', 'key2':'value2'}

```
    """
```

```
s3 = boto.connect_s3()
```

```
bucket = s3.lookup(bucket_name)
```

```
# Get a new, blank Key object from the bucket. This Key object
only
# exists locally until we actually store data in it.
key = bucket.new_key(key_name)
# Add the metadata to the Key object
key.metadata.update(metadata)

# Now, write the data and metadata to S3
key.set_contents_from_filename(path_to_file)
return key
```

```
def print_key_metadata(bucket_name, key_name):  
    """  
    Print the metadata associated with an S3 Key object.  
  
    bucket_name The name of the S3 Bucket.  
    key_name    The name of the object containing the data in  
S3.  
    """  
    s3 = boto.connect_s3()  
    bucket = s3.lookup(bucket_name)
```



```
key = bucket.lookup(key_name)
print key.metadata
```

Computing Total Storage Used by a Bucket

```
import boto
```

```
def bucket_du(bucket_name):
```

```
    """
```

```
        Compute the total bytes used by a bucket.
```

```
        NOTE: This iterates over every key in the bucket. If you have  
millions of
```

```
        keys this could take a while.
```

```
    """
```

```
    s3 = boto.connect_s3()
```

```
    total_bytes = 0
```

```
    bucket = s3.lookup(bucket_name)
```

```
    if bucket:
```

```
        for key in bucket:
```

```
            total_bytes += key.size
```

```
else:  
    print 'Warning: bucket %s was not found!' % bucket_name  
    return total_bytes
```

Copy an Existing Object to Another Bucket

```
import boto
```

```
def copy_object(src_bucket_name,  
                src_key_name,  
                dst_bucket_name,  
                dst_key_name,  
                preserve_metadata=True):
```

```
    """
```

Copy an existing object to another location.

src_bucket_name Bucket containing the existing object.

src_key_name Name of the existing object.

dst_bucket_name Bucket to which the object is being copied.

dst_key_name The name of the new object.

preserve_metadata If True, all metadata on the original object

will be preserved on the new object. If False the new object will have the default metadata.

```
    """
```

```
s3 = boto.connect_s3()
bucket = s3.lookup(src_bucket_name)
# Lookup the existing object in S3
key = bucket.lookup(src_key_name)
# Copy the key back on to itself, with new metadata
return key.copy(dst_bucket_name, dst_key_name,
preserve_acl=preserve_acl)
```

Modify the Metadata of an Existing Object

```
import boto
```

```
def modify_metadata(bucket_name,  
                    key_name,  
                    metadata):
```

```
    """
```

Update the metadata with an existing object.

bucket_name The name of the S3 Bucket.

key_name The name of the object containing the data in S3.

metadata A Python dict object containing the new metadata.

For example: {'key1':'value1', 'key2':'value2'}

```
    """
```

```
s3 = boto.connect_s3()
```

```
bucket = s3.lookup(bucket_name)
```



```
# Lookup the existing object in S3
key = bucket.lookup(key_name)
# Copy the key back on to itself, with new metadata
key.copy(bucket.name, key.name, metadata,
preserve_acl=True)
return key
```

Enable Logging on an Existing Bucket

```
import boto
```

```
def enable_logging(bucket_name,  
                  log_bucket_name,  
                  log_prefix=None):
```

```
    """
```

Enable logging on a bucket.

bucket_name Bucket to be logged.

log_bucket_name Bucket where logs will be written.

log_prefix A string which will be prepended to all log file names.

```
    """
```

```
s3 = boto.connect_s3()
```

```
bucket = s3.lookup(bucket_name)
```

```
log_bucket = s3.lookup(log_bucket_name)
```

```
# First configure log bucket as a log target.
```

```
# This sets permissions on the bucket to allow S3 to write logs.
```

```
log_bucket.set_as_logging_target()
```

```
# Now enable logging on the bucket and tell S3
# where to deliver the logs.
bucket.enable_logging(log_bucket, target_prefix=log_prefix)
def disable_logging(bucket_name):
    """
    Disable logging on a bucket.
    bucket_name    Bucket that will no longer be logged.
    """
```

```
s3 = boto.connect_s3()  
bucket = s3.lookup(bucket_name)  
bucket.disable_logging()
```

Reduce the Cost of Storing Noncritical Data

```
import boto
import os
def upload_file_rrs(local_file,
                    bucket_name,
                    key_name=None):
    """
    Upload a local file to S3 and store is using Reduced
    Redundancy Storage.
    local_file Path to local file.
    bucket_name Bucket to which the file will be uploaded.
    key_name Name of the new object in S3. If not provided,
    the basename
           of the local file will be used.
    """
    s3 = boto.connect_s3()
    bucket = s3.lookup(bucket_name)
    # Expand common shell vars in filename.
    local_file = os.path.expanduser(local_file)
    local_file = os.path.expandvars(local_file)
```



```
# If key_name was not provided, use basename of file.  
if not key_name:  
    key_name = os.path.basename(local_file)  
    # Create a new local key object.  
    key = bucket.new_key(key_name)  
    # Now upload file to S3  
    key.set_contents_from_filename(local_file,  
    reduced_redundancy=True)
```

```
def copy_object_to_rrs(bucket_name,
                       key_name):
    """
    Will change an existing standard storage class object to a
    Reduced Redundancy storage class object.

    bucket_name Bucket in which the existing key is located.
    key_name    Name of the existing, standard storage key.
    """
    s3 = boto.connect_s3()
    bucket = s3.lookup(bucket_name)
    key = bucket.lookup(key_name)

    return key.copy(bucket_name, key_name,
                    reduced_redundancy=True,
                    preserve_acl=True)
```

Hosting Static Websites on S3

```
import boto
import os
import time
def upload_website(bucket_name,
                  website_dir,
                  index_file,
                  error_file=None):
```

```
    """
```

Upload a static website contained in a local directory to a bucket in S3.

bucket_name The name of the bucket to upload website to.

website_dir Fully-qualified path to directory containing website.

index_file The name of the index file (e.g. index.html)

error_file The name of the error file. If not provided the default S3 error page will be used.

```
    """
```

```
s3 = boto.connect_s3()
bucket = s3.lookup(bucket_name)
# Make sure bucket is publicly readable
bucket.set_canned_acl('public-read')
for root, dirs, files in os.walk(website_dir):
    for file in files:
        full_path = os.path.join(root, file)
        rel_path = os.path.relpath(full_path, website_dir)
        print 'Uploading %s as %s' % (full_path, rel_path)
        key = bucket.new_key(rel_path)
        key.content_type = 'text/html'
        key.set_contents_from_filename(full_path,
policy='public-read')
# Now configure the website
bucket.configure_website(index_file, error_file)
# A short delay, just to let things become consistent.
time.sleep(5)
print 'You can access your website at:'
print bucket.get_website_endpoint()
```